

Image Processing using Graph Laplacian Operator

David Wobrock

ALPINES Team - INRIA Paris
Laboratoire Jacques-Louis Lions - Sorbonne Université
KTH, Stockholm
INSA Lyon

May 18, 2018

Table of Contents

Introduction

Image processing using Laplacian operator

Parallel implementation

Conclusion

Background

- ▶ Millions of pictures shared daily
- ▶ Image processing using spectral methods on smartphones (denoising, sharpening, ...)
- ▶ Involves eigenvalue problems, linear algebra and solving dense linear systems
- ▶ Opportunity for high-performance computing and parallelism

Objective

- ▶ Not necessarily improving image processing
- ▶ Analyse the behaviour of solving large dense systems
- ▶ Large: N^2 , N the number of pixels in the input image
- ▶ Dense: affinity and Laplacian matrices are dense because global filter

Image processing - Global filter algorithm

- ▶ z output image, y input image, W data-dependent global filter of size N^2
- ▶ $z_i = \sum_{j=1}^N W_{ij} y_j$
- ▶ A vector of weights for each pixel
- ▶ Global filter: $z = Wy$

Image processing - Image as graph

- ▶ Think of the image as a graph
- ▶ Each pixel is a node
- ▶ The graph is complete and the edges are weighted
- ▶ The weight represents the similarity between two pixels/nodes
- ▶ Similarity can be measured by weighting together spatial and color closeness of pixels
- ▶ Bilateral kernel: $\exp\left(-\frac{\|x_i - x_j\|^2}{h_x^2}\right) \cdot \exp\left(-\frac{\|z_i - z_j\|^2}{h_z^2}\right)$

Image processing - Filter and Laplacian

- ▶ K affinity matrix, \mathcal{L} Laplacian matrix, W filter matrix and f a function
- ▶ Filter defined as $W = I - f(\mathcal{L})$
- ▶ Let $d_i = \sum_j K_{ij}$, $D = \text{diag}\{d_i\}$ and $\bar{d} = \frac{1}{N} \sum_i d_i$

Laplacian	Formula of \mathcal{L}	Symmetric	Spectral Range
Un-normalised	$D - K$	Yes	$[0, n]$
Normalised	$I - D^{-1/2}KD^{-1/2}$	Yes	$[0, 2]$
Random walk	$I - D^{-1}K$	No	$[0, 1]$
"Sinkhorn"	$I - C^{-1/2}KC^{-1/2}$	Yes	$[0, 1]$
Re-normalised	$\alpha(D - K)$, $\alpha \approx \bar{d}^{-1}$	Yes	$[0, n]$

Table: Overview of different graph Laplacian operator definitions.

Image processing - Sampling and Nyström extension

- ▶ Sample p pixels of the image ($\leq 1\%$ of pixels)
- ▶ With $K_A \in \mathbb{R}^{p \times p}$, $K_B \in \mathbb{R}^{p \times N-p}$ and $K_C \in \mathbb{R}^{N-p \times N-p}$
$$K = \begin{bmatrix} K_A & K_B \\ K_B^T & K_C \end{bmatrix}$$
- ▶ Approximate K by $\tilde{K} = \tilde{\Phi} \tilde{\Pi} \tilde{\Phi}^T$
- ▶ With submatrix $K_A = \Phi_A \Pi_A \Phi_A^T$,
$$\tilde{\Pi} = \Pi_A$$
$$\tilde{\Phi} = \begin{bmatrix} \Phi_A \\ K_B^T \Phi_A \Pi_A^{-1} \end{bmatrix}$$
- ▶ Finally, $\tilde{K} = \begin{bmatrix} K_A & K_B \\ K_B^T & K_B^T K_A^{-1} K_B \end{bmatrix}$

Image processing - Eigenvalues

- ▶ For the filter, only the largest eigenvalues are needed because the smallest tends to 0
- ▶ There is an equivalence between largest eigenvalues of W and the smallest ones of \mathcal{L}
- ▶ The smallest eigenvalues can be computed with the inverse power method, which requires solving systems of linear equations
- ▶ But the Nyström extension only works for leading eigenvalues.

Parallel implementation details

- ▶ C language
- ▶ PETSc (Toolkit for scientific computing in parallel)
- ▶ SLEPc, Elemental

Full matrix computation result



Figure: Left: input image. Right: sharpened image.

Runtime of full matrix computation

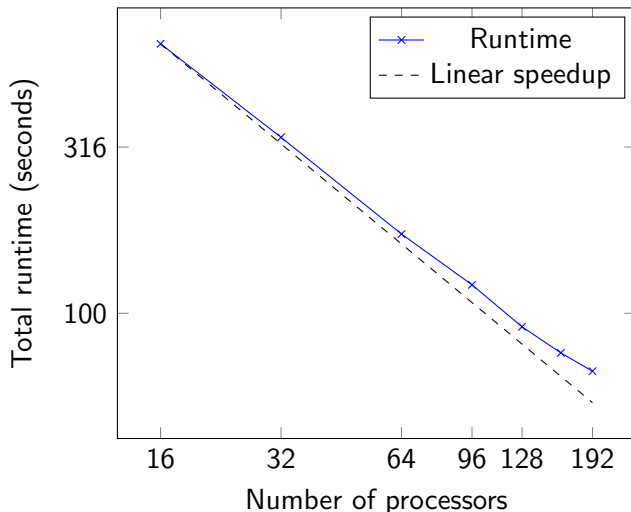


Figure: Total runtime of the algorithm with entire matrix computation (log scale).

Inverse subspace iteration

Algorithm 2 Inverse subspace iteration

Input: A the matrix of size $p \times p$, m the number of required eigenvalues, ε a tolerance

Output: X_k the desired invariant subspace

Initialise m random orthonormal vectors X_0 of size p

$$R_0 \leftarrow (I - X_0 X_0^T) A X_0$$

For $k=0, 1, 2, \dots$

while $\|R_k\| > \varepsilon$ **do**

for $i=1$ **to** m **do**

$$\text{Solve } A X_{k+1}^{(i)} = X_k^{(i)}$$

end for


$$X_{k+1} \leftarrow \text{Orthonormalise}(X_{k+1})$$

$$R_{k+1} \leftarrow (I - X_{k+1} X_{k+1}^T) A X_{k+1}$$

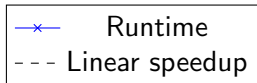

end while

Approximation runtime of approximated computation

50 eigenvalues.



500 eigenvalues.



Runtime in the inverse subspace iteration algorithm

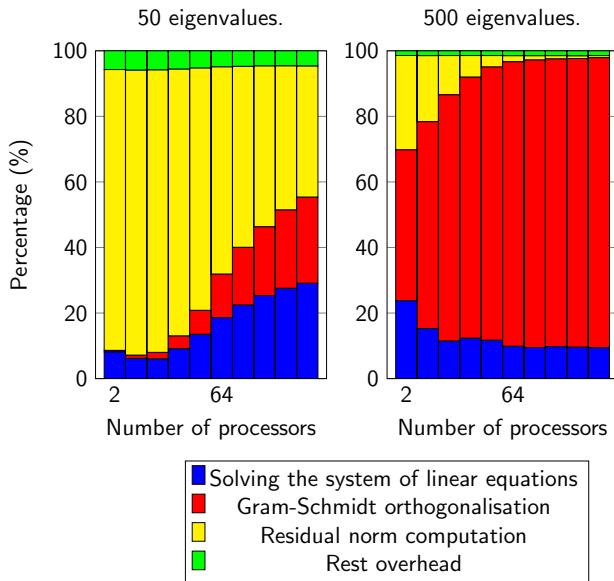



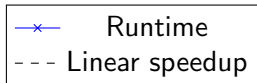

Figure: Proportion of each step in the inverse subspace iteration.

Skipping the Gram-Schmidt procedure

50 eigenvalues.



500 eigenvalues.



Skipping Gram-Schmidt more often

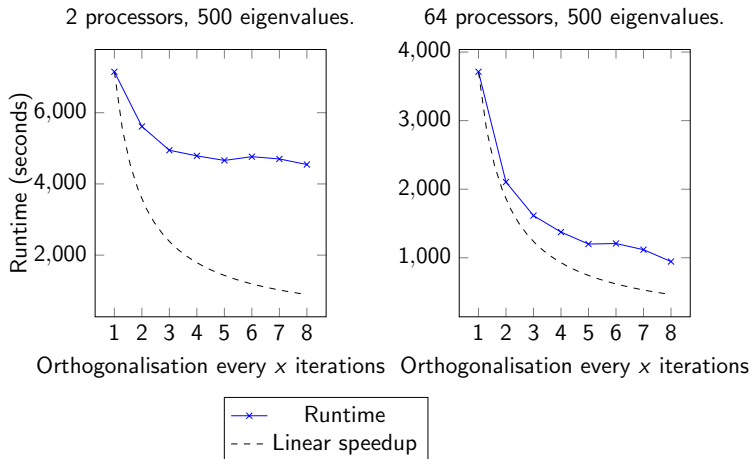


Figure: Runtime of the inverse subspace iteration depending on the amount of Gram-Schmidt procedures.

Linear solver performances




Figure: Runtime of the linear solver for 50 eigenvalues for a 4000×4000 matrix (log scale).

Linear solver performances - big image




Figure: Runtime of the linear solver for 50 eigenvalues for a 14400×14400 matrix (log scale).

Discussions & perspectives

- ▶ Linear solver with domain decomposition methods as preconditioner
- ▶ Skipping Gram-Schmidt
- ▶ Finished image processing algorithm
- ▶ State-of-the-art performances with a similar method
- ▶ Explore more solvers, domain decomposition methods and other preconditioner